

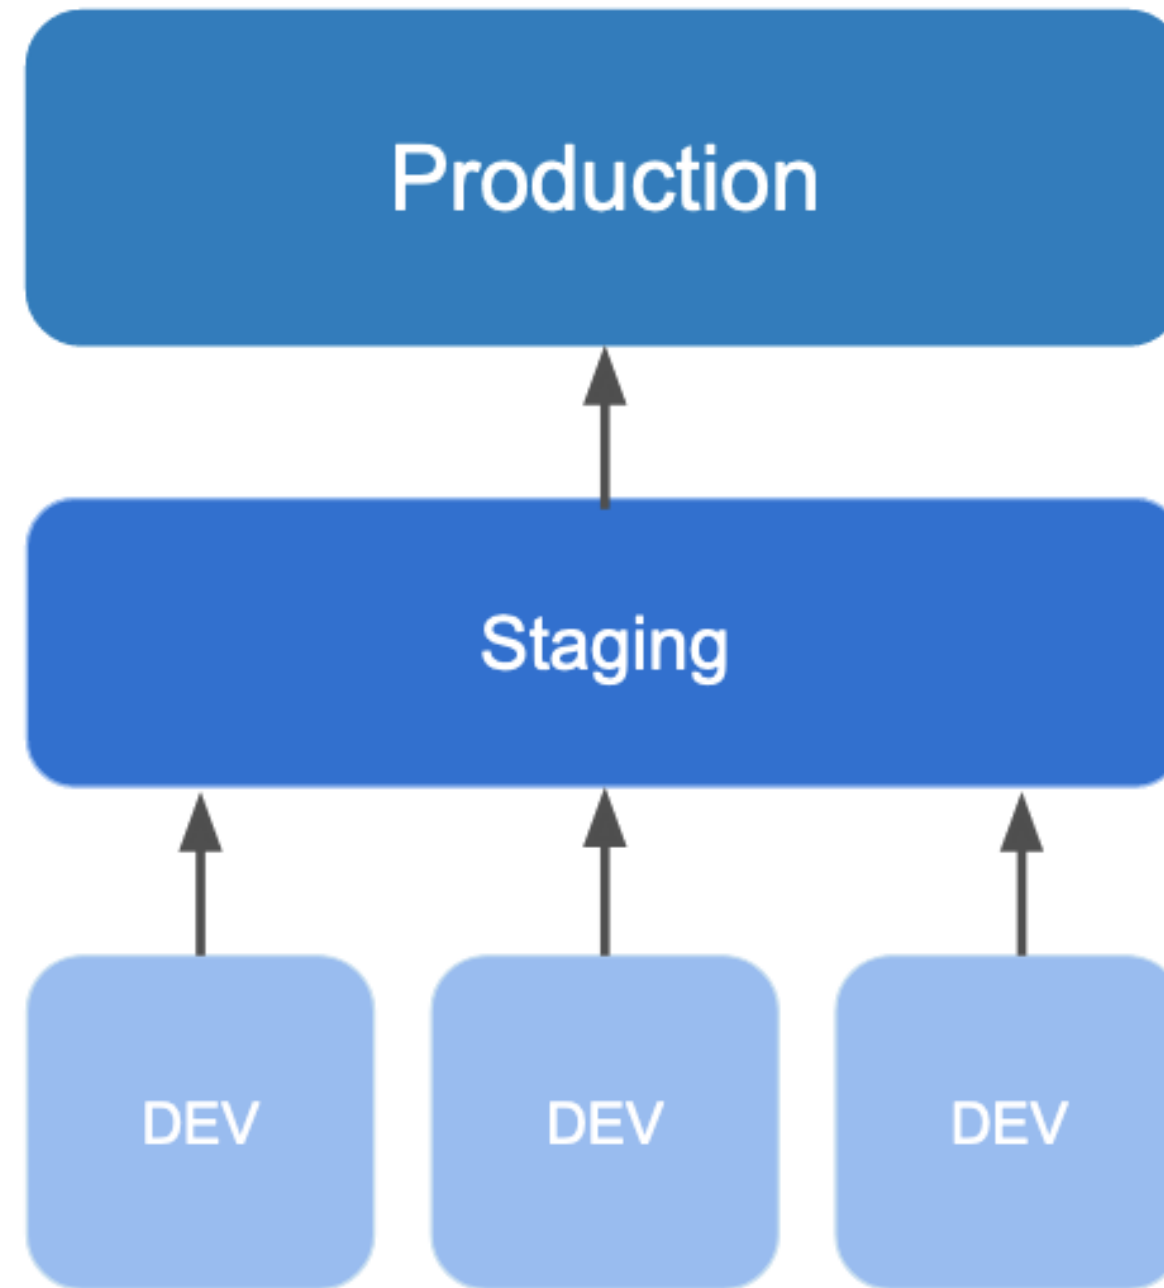
Introduction to MLflow Model Registry

INTRODUCTION TO MLFLOW



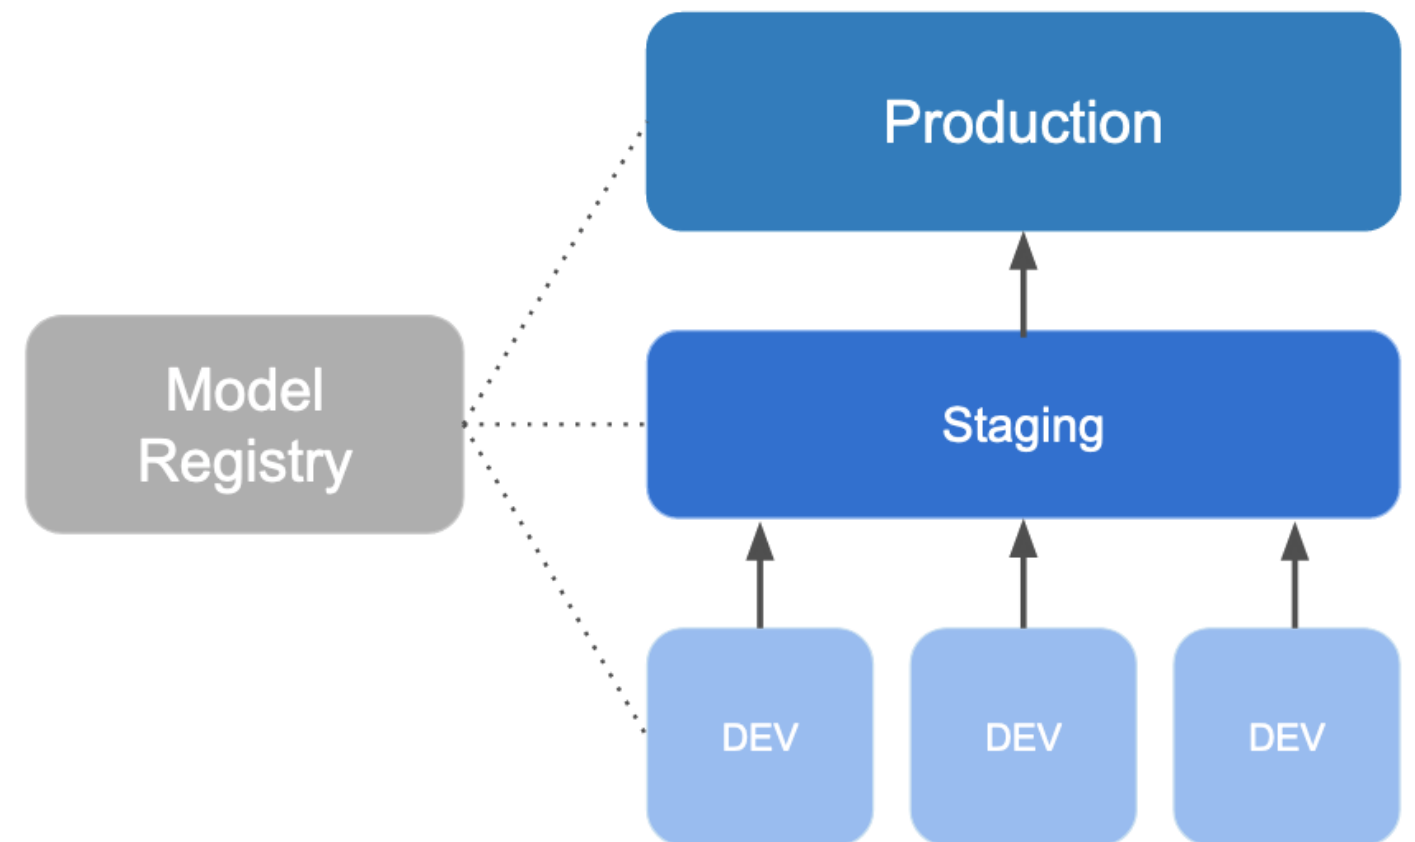
Weston Bassler
Senior MLOps Engineer

Model lifecycle



MLflow Model Registry

- Centralized storage location
- Lifecycle management
 - Web UI
 - MLflow Client module
 - Model versions
 - Model stages

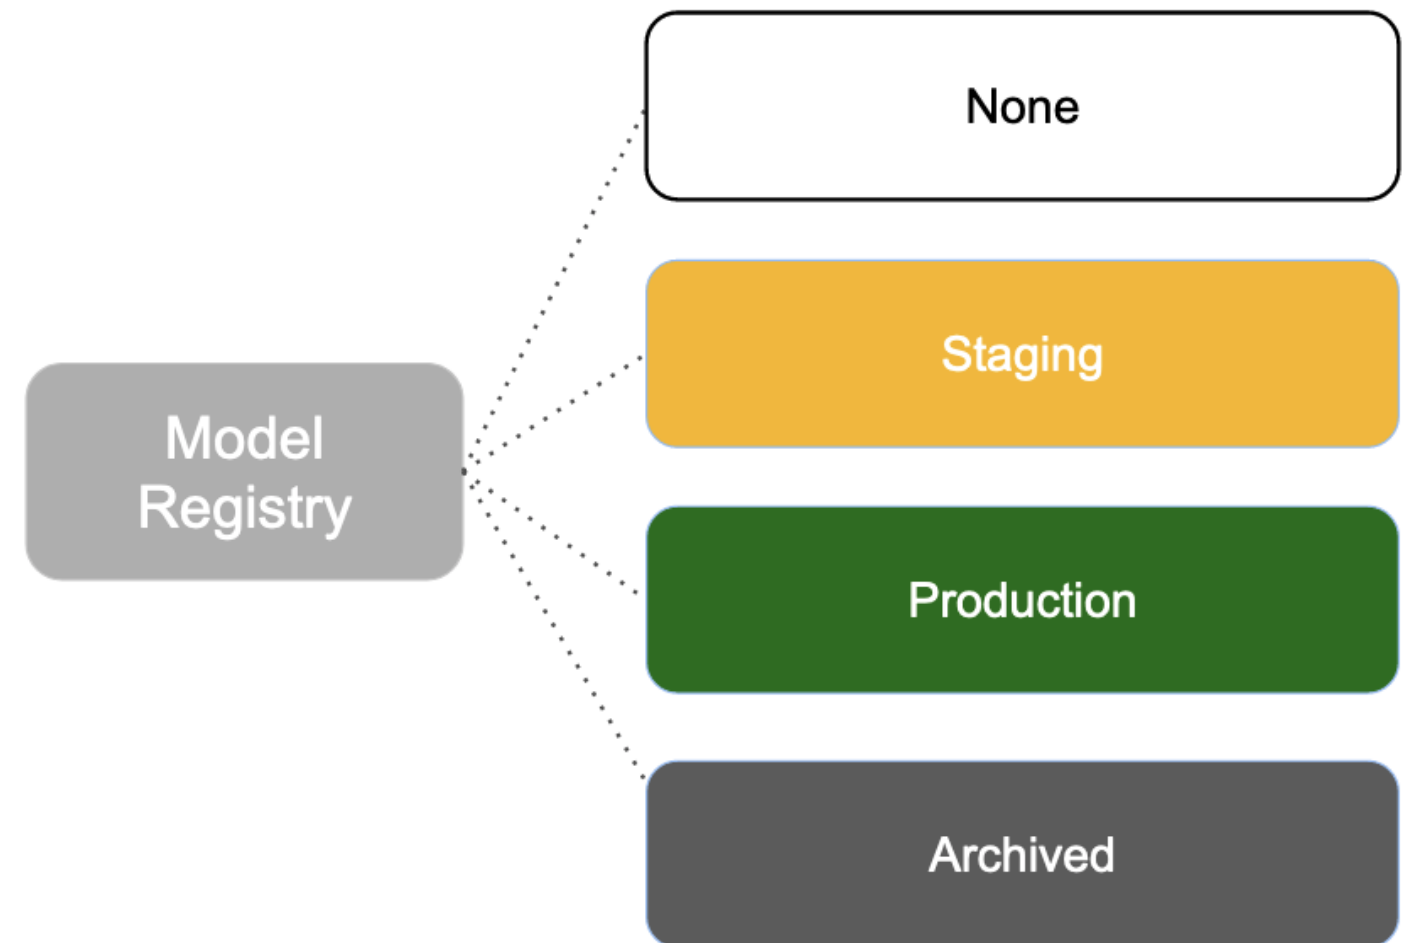


MLflow Model Registry



- Model
 - Model logged to MLflow Tracking
- Registered Model
 - Obtains version
 - Stage eligible

MLflow Model Registry

- Model Version
 - Increments with each new registered model
- Model Stage
 - Can be assigned one of:
 - None
 - Staging
 - Production
 - Archived





Working with the Model Registry




Models

Registered Models

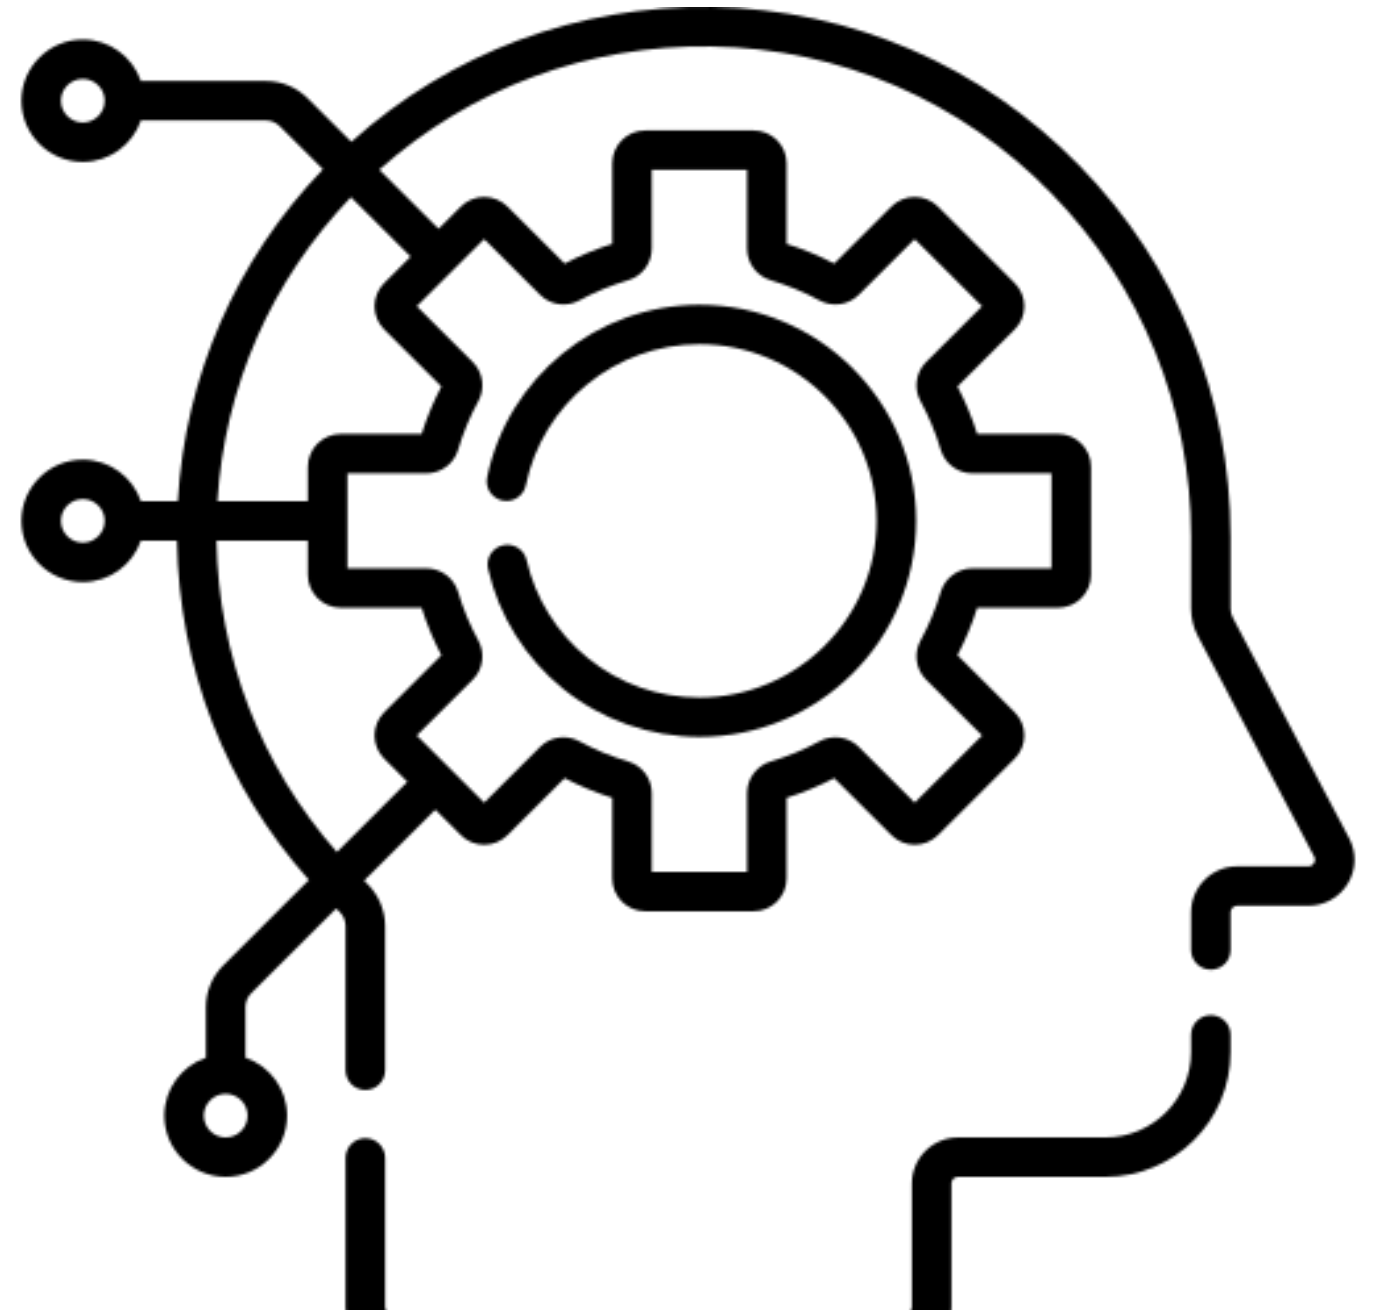
Create Model

  Search by model names or tags

Name 	Latest Version	Staging	Production	Last Modified
No models yet. Create a model to get started.				

MLflow Client module

- Experiments
- Runs
- Model Versions
- Registered Models



¹ Flaticon.com

Using MLflow client module

```
# Import from MLflow module
from mlflow import MlflowClient

# Create an instance
client = MlflowClient()

# Print the object
client
```

```
<mlflow.tracking.client.MlflowClient object at 0x101d55f30>
```


Registering a model

```
# Create a Model named "Unicorn"  
client.create_registered_model(name="Unicorn")
```

```
<RegisteredModel: creation_timestamp=1679404160448, description=None,  
last_updated_timestamp=1679404160448, latest_versions=[], name='Unicorn',  
tags={}>
```

Model UI

Registered Models

Create Model

?

Q

Search by model names or tags

Search

Clear

Name	Latest Version	Staging	Production	Last Modified	Tags
Unicorn	—	—	—	2023-03-21 09:09:20	—

Searching registered models

Registered Models

Create Model

?

Q Search by model names or tags

Search

Clear

Name	Latest Version	Staging	Production	Last Modified	Tags
Unicorn	—	—	—	2023-03-21 09:09:20	—

Searching registered models

```
# Search for registered models
```

```
client.search_registered_models(filter_string=MY_FILTER_STRING)
```

- Identifiers
 - name - of the model
 - tags - tags associated with model
- Comparators
 - `=` - equal to
 - `!=` - not equal to
 - `LIKE` - case-sensitive pattern match
 - `ILIKE` - case-insensitive pattern match

Example search

```
# Filter string
```

```
unicorn_filter_string = "name LIKE 'Unicorn%'"
```

```
# Search models
```

```
client.search_registered_models(filter_string=unicorn_filter_string)
```

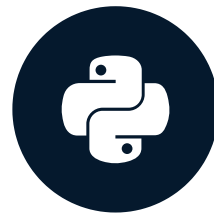
```
[<RegisteredModel: creation_timestamp=1679404160448, description=None,  
last_updated_timestamp=1679404160448, latest_versions=[], name='Unicorn',  
tags={}>,  
<RegisteredModel: creation_timestamp=1679404276745, description=None,  
last_updated_timestamp=1679404276745, latest_versions=[], name='Unicorn 2.0',  
tags={}>]
```

Let's practice!

INTRODUCTION TO MLFLOW

Registering Models

INTRODUCTION TO MLFLOW



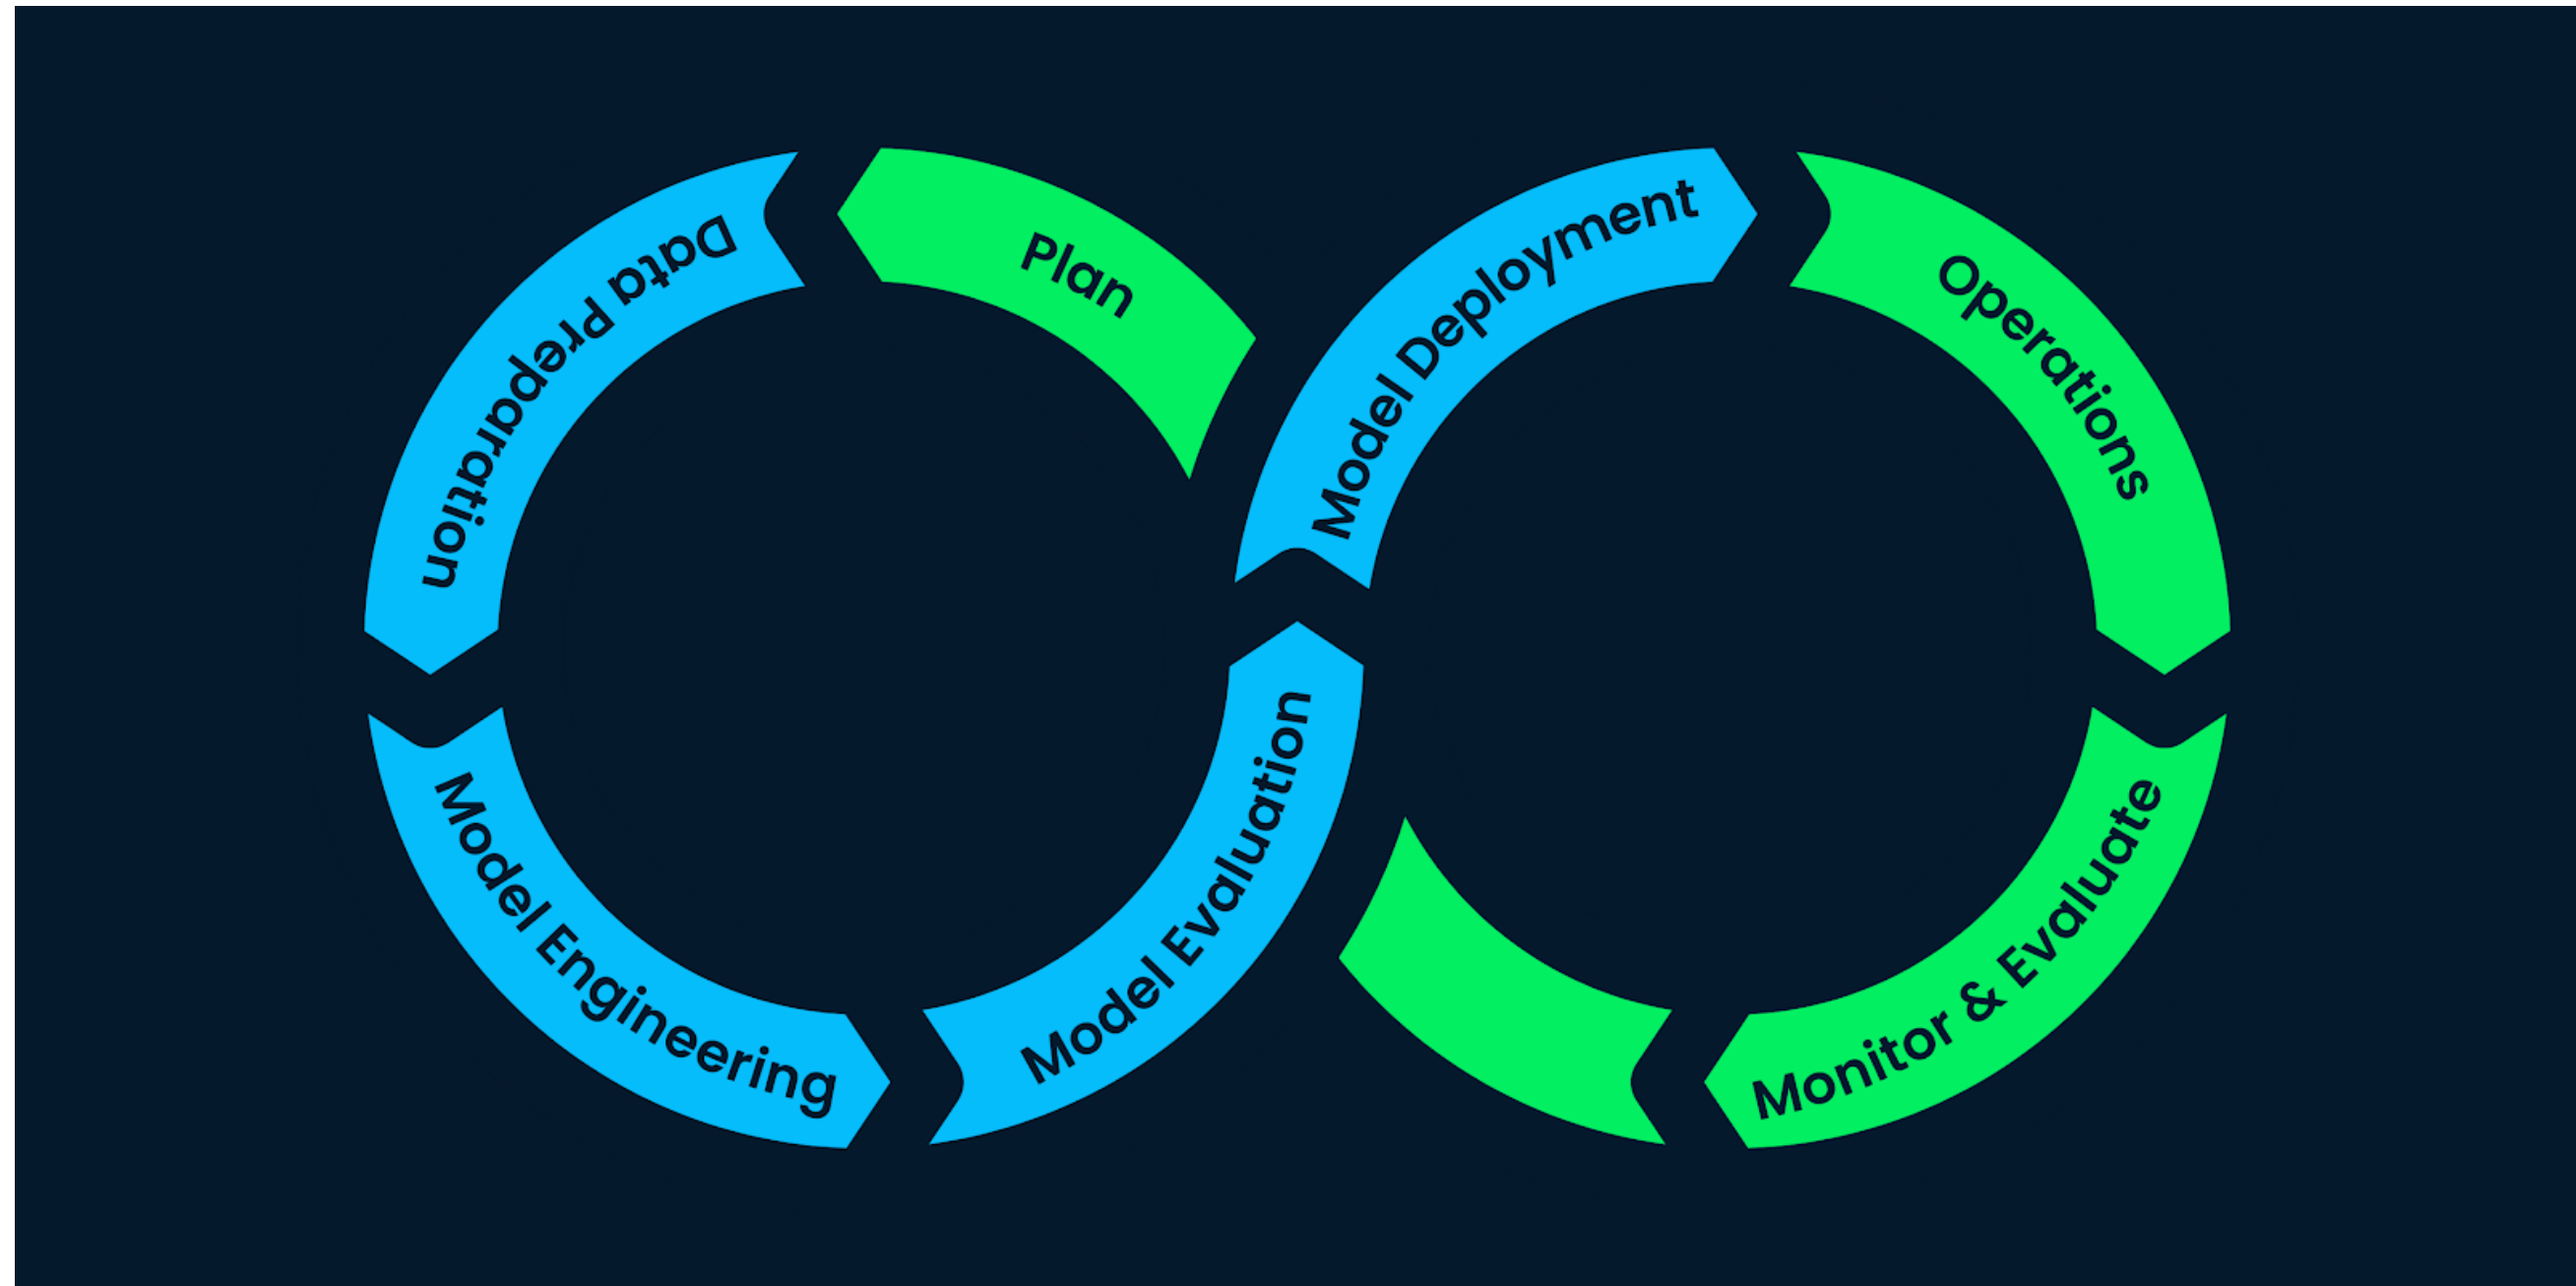
Weston Bassler

Senior MLOps Engineer

Registering MLflow Models

- Model Versions
 - follows traditional software development
 - track changes
- Collaboration
 - between different roles
 - same roles for improvement

Model lifecycle management



¹ datacamp.com

Ways to register models

```
# Existing MLflow Models
```

```
mlflow.register_model(model_uri, name)
```

model_uri

- local filesystem
- tracking server

```
# During training run
```

```
mlflow.FLAVOR.log_model(name,  
    artifact_uri,  
    registered_model_name="MODEL_NAME")
```

registered_model_name="MODEL_NAME"

Registering model example

```
# Import mlflow
import mlflow

# Register model from local filesystem
mlflow.register_model("./model", "Unicorn")

# Register model from Tracking server
mlflow.register_model("runs:/run-id/model", "Unicorn")
```

```
# Register local MLFlow Model
```

```
mlflow.register_model(model_uri="./model", name="Unicorn")
```

```
Registered model 'Unicorn' already exists. Creating a new version of this model...
```

```
2023/03/24 14:34:26 INFO mlflow.tracking._model_registry.client:
```

```
Waiting up to 300 seconds for model version to finish creation.
```

```
Model name: Unicorn, version 1
```

```
Created version '1' of model 'Unicorn'.
```

```
<ModelVersion: creation_timestamp=1679682866413, current_stage='None',  
description=None, last_updated_timestamp=1679682866413, name='Unicorn',  
run_id=None, run_link=None, source='./model', status='READY', status_message=None,  
tags={}, user_id=None, version=1>
```

```
# Register model from MLflow Tracking
```

```
mlflow.register_model(model_uri="runs:/run-id/model", name="Unicorn")
```

```
Registered model 'Unicorn' already exists. Creating a new version of this model...
```

```
2023/03/24 14:36:56 INFO mlflow.tracking._model_registry.client:
```

```
Waiting up to 300 seconds for model version to finish creation.
```

```
Model name: Unicorn, version 2
```

```
Created version '2' of model 'Unicorn'.
```

```
<ModelVersion: creation_timestamp=1679683016297, current_stage='None',  
description=None, last_updated_timestamp=1679683016297, name='Unicorn',  
run_id='2e974508b68b45ceb114657c6e97fef5', run_link=None,  
source='./mlruns/1/2e974508b68b45ceb114657c6e97fef5/artifacts/model',  
status='READY', status_message=None, tags={}, user_id=None, version=2>
```

Models UI

Registered Models

Create Model

Name	Latest Version
Insurance	—
Insurance2	—
Test Scores	—
Unicorn	Version 2
Unicorn 2.0	—

Unicorn versions

[Registered Models](#) >

Unicorn

Created Time: 2023-03-21 09:09:20

> Description [Edit](#)

> Tags

▼ Versions

All

Active 0

Compare

<input type="checkbox"/>	Version	Registered at
<input type="checkbox"/>	<input checked="" type="checkbox"/> Version 2	2023-03-24 14:36:56
<input type="checkbox"/>	<input checked="" type="checkbox"/> Version 1	2023-03-24 14:34:26

Logging model

```
# Import modules
import mlflow
import mlflow.sklearn
from sklearn.linear_model import LogisticRegression

# Model
lr = LogisticRegression()
lr.fit(X, y)

# Log model
mlflow.sklearn.log_model(lr, "model", registered_model_name="Unicorn")
```



```
# Log model
```

```
mlflow.sklearn.log_model(lr, "model", registered_model_name="Unicorn")
```

```
Registered model 'Unicorn' already exists. Creating a new version of this model...
```

```
2023/03/24 17:31:10 INFO mlflow.tracking._model_registry.client:
```

```
Waiting up to 300 seconds for model version to finish creation.
```

```
Model name: Unicorn, version 3
```

```
Created version '3' of model 'Unicorn'.
```

```
<mlflow.models.model.ModelInfo object at 0x14734d330>
```

Let's practice

INTRODUCTION TO MLFLOW

Model stages

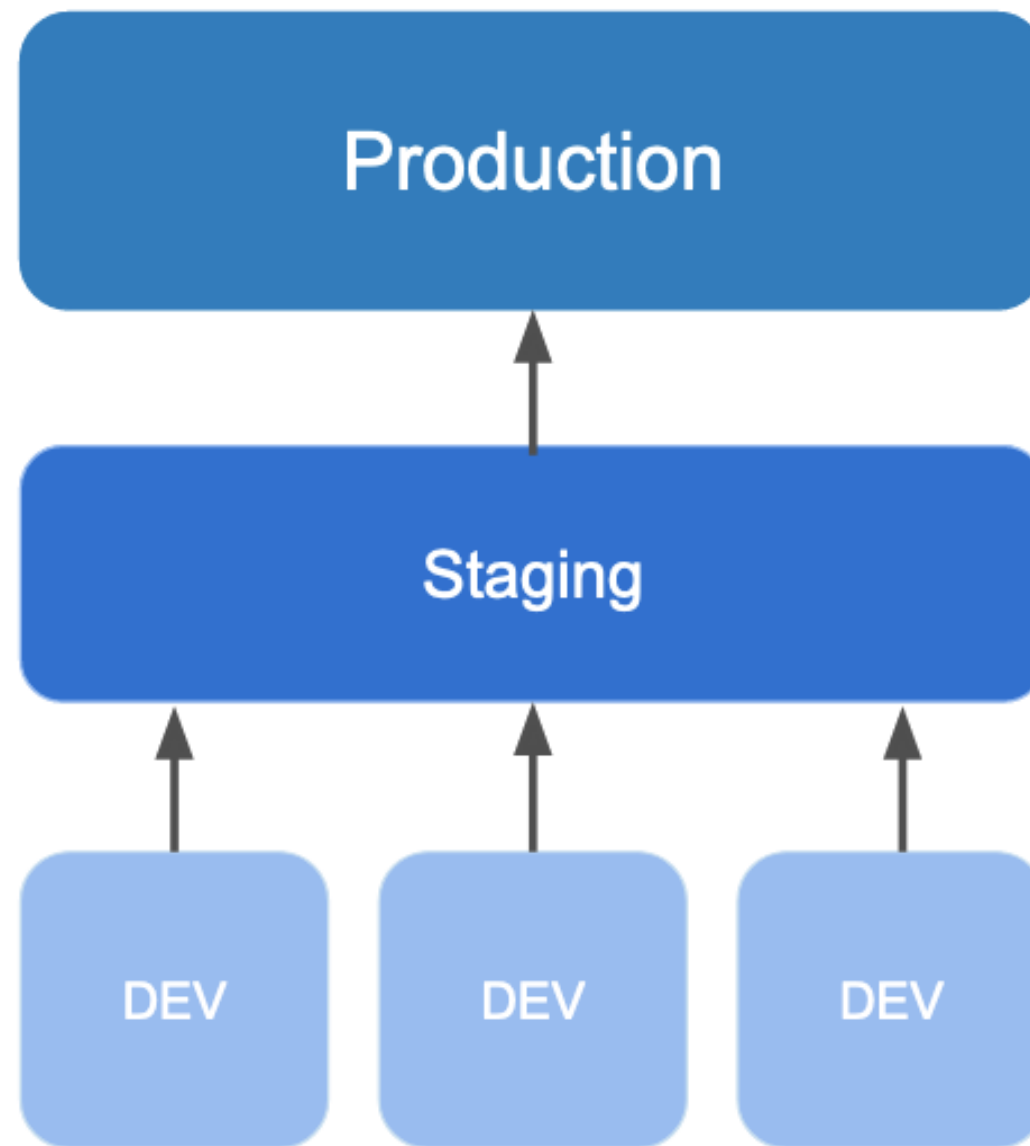
INTRODUCTION TO MLFLOW



Weston Bassler

Senior MLOps Engineer

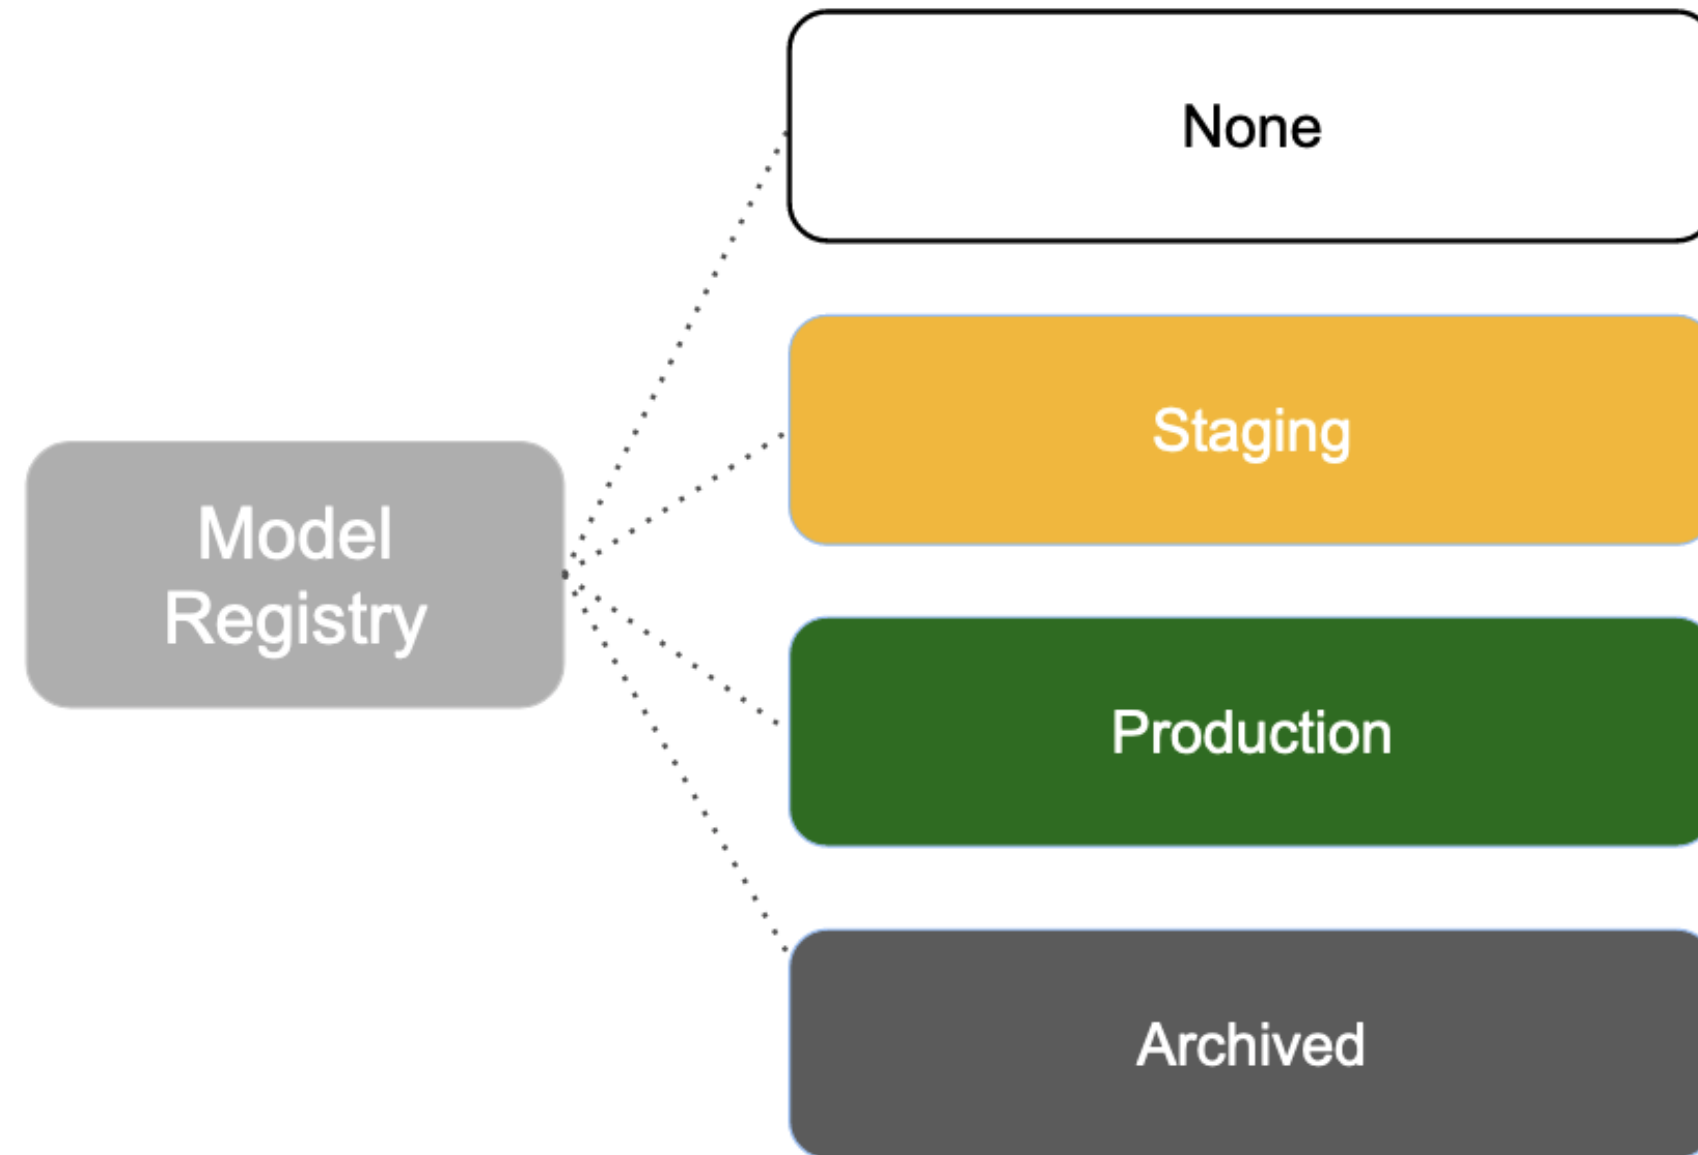
Software environments



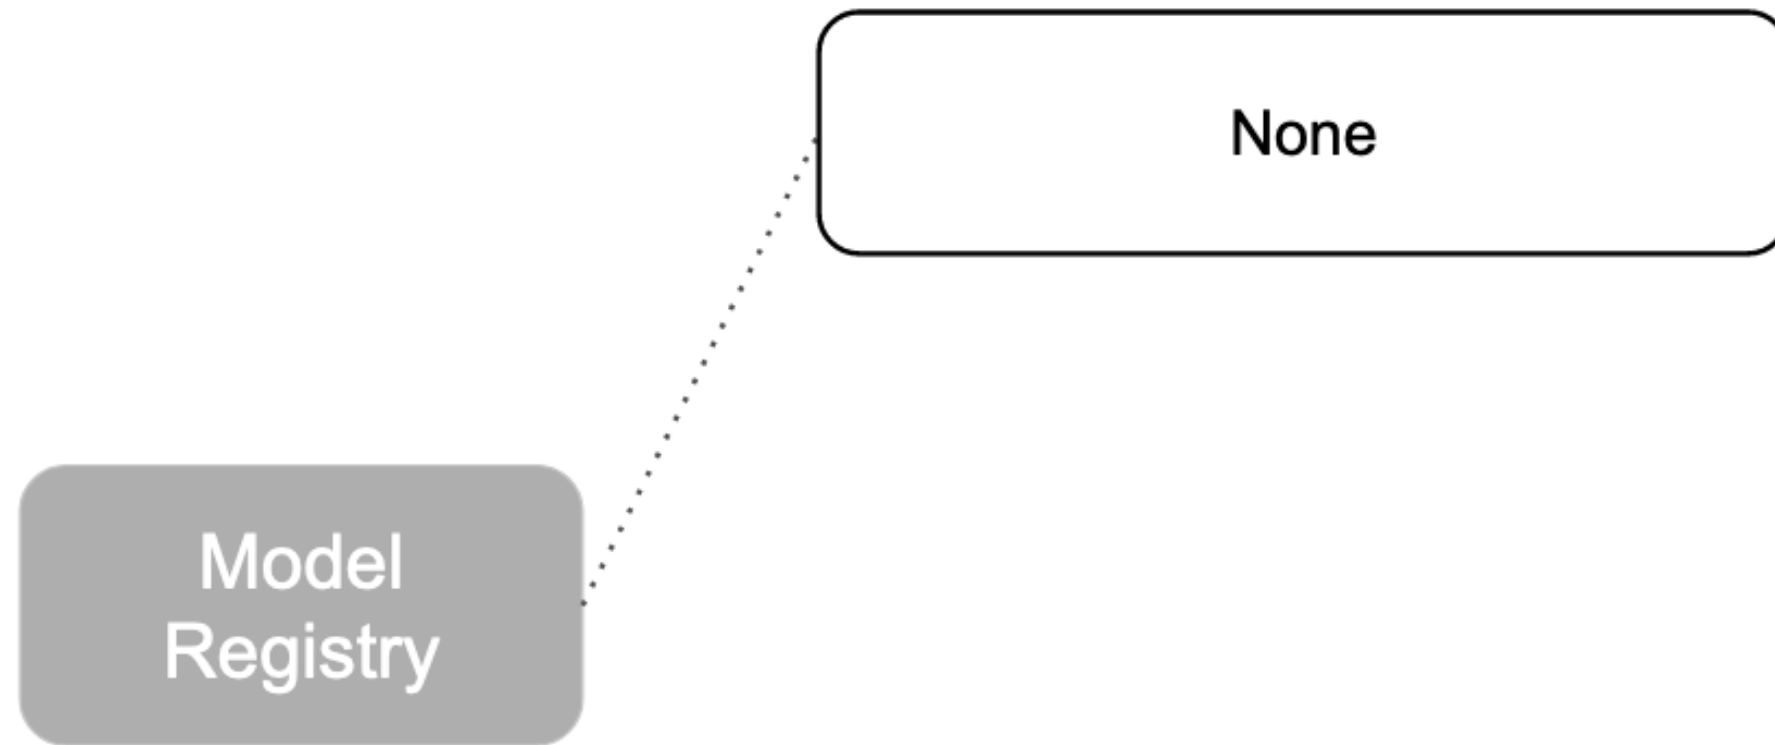
MLflow model stages

- Assigned to model versions
- Predefined stages:
 - None
 - Staging
 - Production
 - Archived
- One single stage at a time

Predefined stages



None



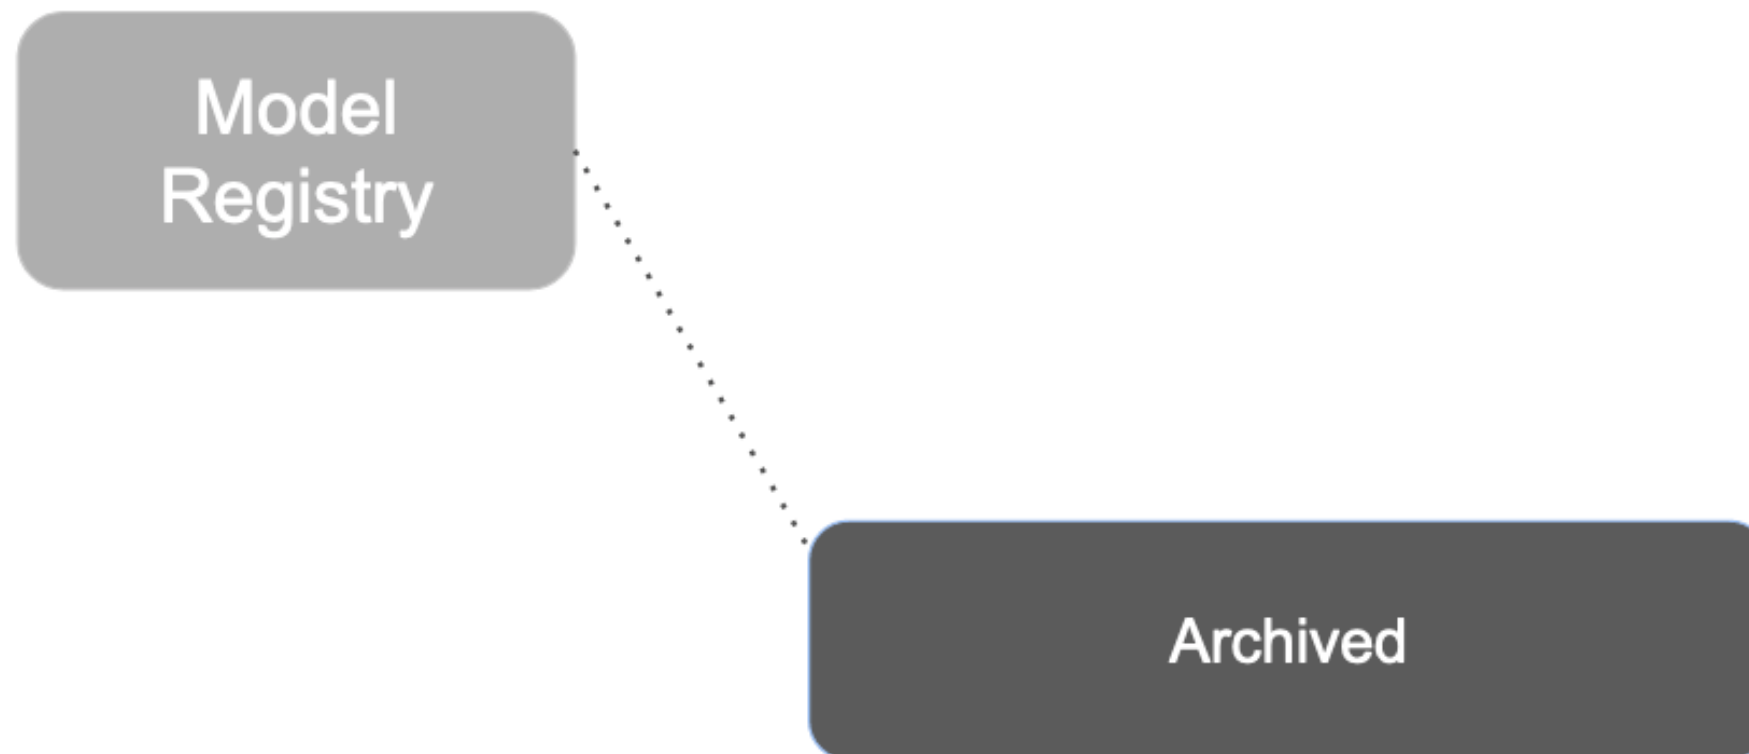
Staging



Production



Archived



Transitioning models

Registered Models > Unicorn >

Version 3

Registered At: 2023-03-24 17:31:10

Stage: **None** ▾

Last Mo Transition to → **Staging**

Source Transition to → **Production**

> D Transition to → **Archived**

```
# Import MLflow Client
from mlflow import MlflowClient
client = MlflowClient()

# Transition to Staging
client.transition_model_version_stage(
    name="Unicorn",
    version=3,
    stage="Staging"
)
```

Transition model version staging

```
# Transition to Staging
```

```
client.transition_model_version_stage(name="Unicorn", version=3, stage="Staging")
```

```
<ModelVersion: creation_timestamp=1679693470034, current_stage='Staging',  
description=None, last_updated_timestamp=1679699050734, name='Unicorn',  
run_id='a1454f2865e449f8835f38f71e53e547', run_link=None,  
source='./mlruns/1/a1454f2865e449f8835f38f71e53e547/artifacts/model',  
status='READY', status_message=None, tags={}, user_id=None, version=3>
```

Registry UI

[Registered Models](#) > [Unicorn](#) >

Version 3

Registered At: 2023-03-24 17:31:10

Stage: **Staging** ▼

Last Modified: 2023-03-24 19:04:10

Source Run: [masked-perch-66](#)

Transitioning to production

```
# Transition to Production
```

```
client.transition_model_version_stage(name="Unicorn", version=3,  
                                     stage="Production")
```

```
<ModelVersion: creation_timestamp=1679693470034, current_stage='Production',  
description=None, last_updated_timestamp=1679699633297, name='Unicorn',  
run_id='a1454f2865e449f8835f38f71e53e547', run_link=None,  
source='./mlruns/1/a1454f2865e449f8835f38f71e53e547/artifacts/model',  
status='READY', status_message=None, tags={}, user_id=None, version=3>
```

Let's practice!

INTRODUCTION TO MLFLOW

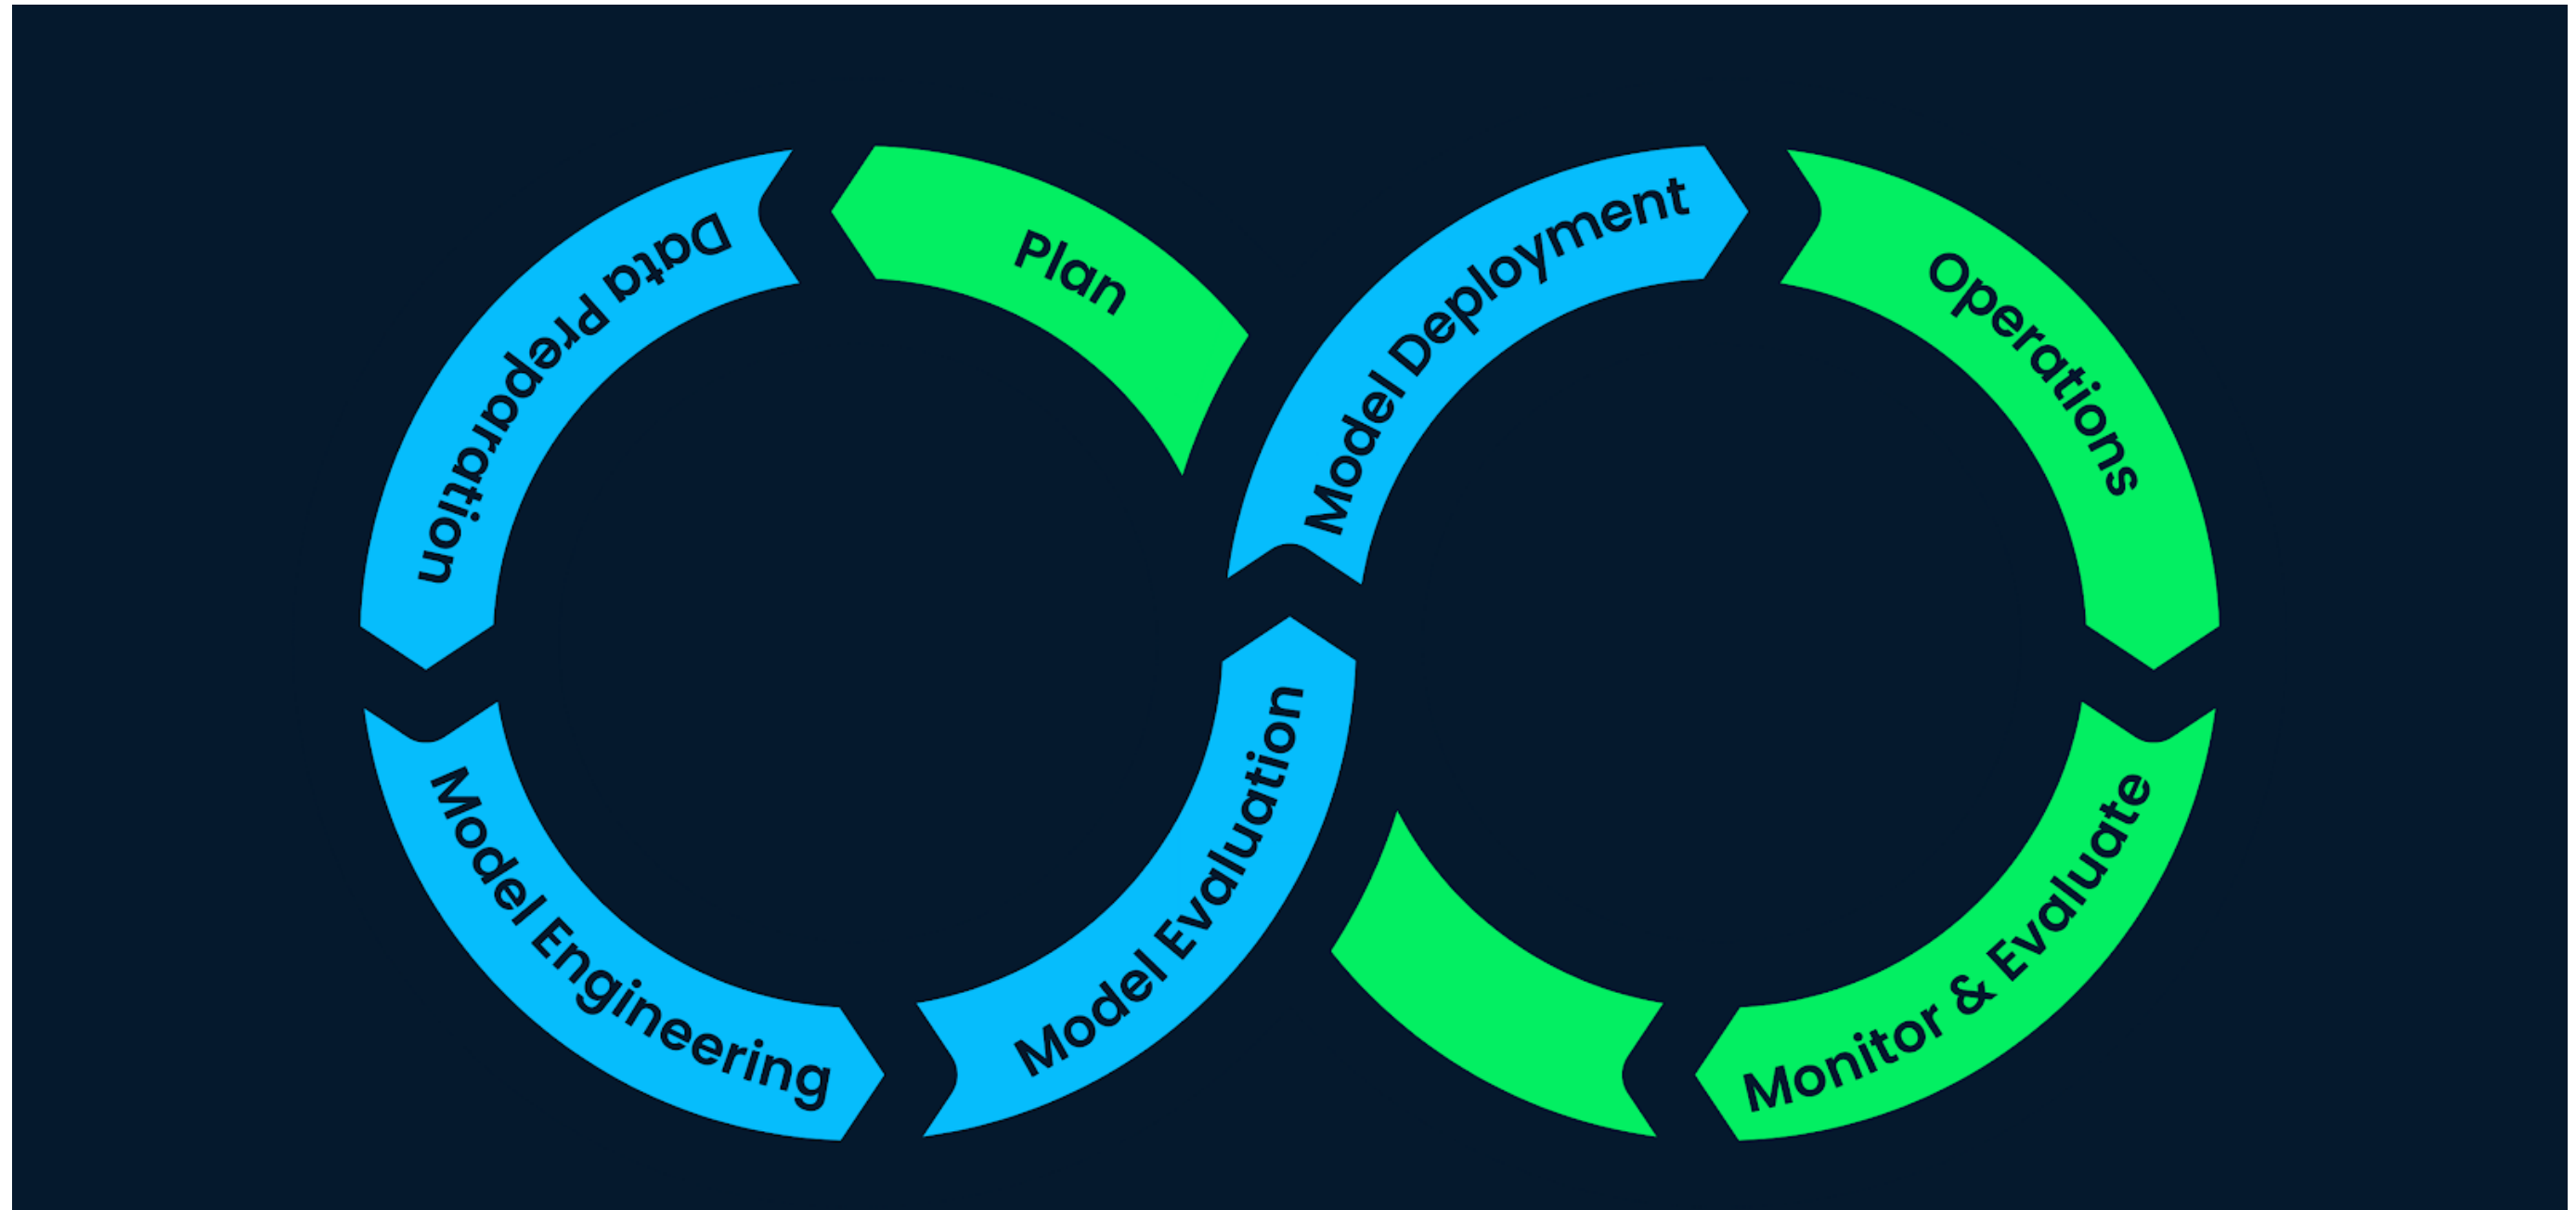
Model deployment

INTRODUCTION TO MLFLOW



Weston Bassler
Senior MLOps Engineer

ML lifecycle



¹ datacamp.com

Model versions and stages

- Model versions
 - Version +1
- Model stages
 - Staging
 - Production
 - Archived

Registered Models

Create Model



Search by model names or tags

Name	Latest Version	Staging	Production
Insurance	Version 1	Version 1	—
Insurance2	Version 2	—	Version 2
Test Scores	—	—	—
Unicorn	Version 3	Version 3	Version 1

Ways to deploy models

Load model

```
# MLflow flavor  
mlflow.FLAVOR.load_model()
```

Serve model

```
# MLflow serve command-line  
mlflow models serve
```

Models URI

Convention

```
models:/
```

Model version

```
models:/model_name/version
```

Model stage

```
models:/model_name/stage
```

Load models

```
# Import flavor
import mlflow.FLAVOR

# Load version
mlflow.FLAVOR.load_model("models:/model_name/version")

# Load stage
mlflow.FLAVOR.load_model("models:/model_name/stage")
```

Load models example

```
# Import flavor
import mlflow.sklearn

# Load Unicorn model in Staging
model = mlflow.sklearn.load_model("models:/Unicorn/Staging")
# Print model
model
```

```
LogisticRegression()
```

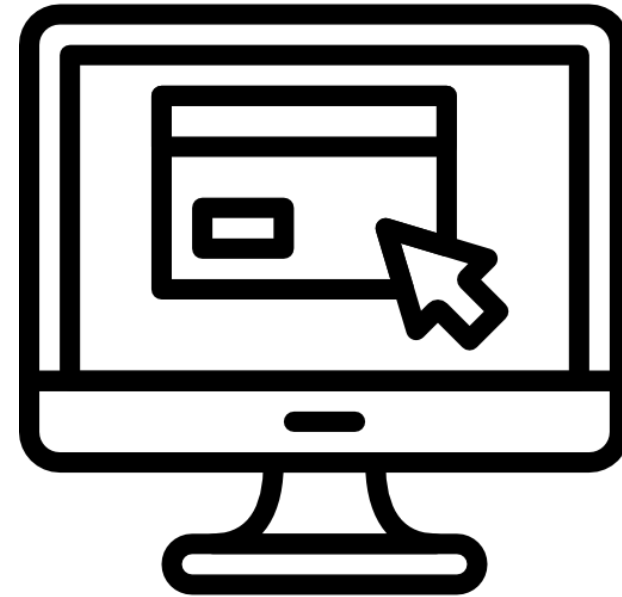
```
# Inference
model.predict(data)
```

Serving models

```
# Serve Unicorn model in Production stage
mlflow models serve -m "models:/Unicorn/Production"
```

```
2023/03/26 15:07:00 INFO mlflow.models.flavor_backend_registry:
Selected backend for flavor 'python_function'
2023/03/26 15:07:00 INFO mlflow.pyfunc.backend: === Running command 'exec gunicorn
--timeout=60 -b 127.0.0.1:5000 -w 1 ${GUNICORN_CMD_ARGS} --
mlflow.pyfunc.scoring_server.wsgi:app'
[2023-03-26 15:07:00 -0400] [86409] [INFO] Starting gunicorn 20.1.0
[2023-03-26 15:07:00 -0400] [86409] [INFO] Listening at: http://127.0.0.1:5000
[2023-03-26 15:07:00 -0400] [86409] [INFO] Using worker: sync
[2023-03-26 15:07:00 -0400] [86410] [INFO] Booting worker with pid: 86410
```

Invocations endpoint



<http://localhost:5000/invocations>

- Formats
 - CSV
 - JSON

¹ Flaticon.com

CSV format

```
pandas_df.to_csv()
```

JSON format

```
{  
  "dataframe_split": {  
    "columns": ["R&D Spend", "Administration", "Marketing Spend", "State"],  
    "data": [["165349.20", 136897.80, 471784.10, 1]]  
  }  
}
```

Model prediction

```
# Send payload to invocations endpoint
curl http://127.0.0.1:5000/invocations -H 'Content-Type: application/json' -d
{
  "dataframe_split": {
    "columns": ["R&D Spend", "Administration", "Marketing Spend", "State"],
    "data": [["165349.20", 136897.80, 471784.10, 1]]
  }
}
```

```
[[104055.1842384]]
```

Let's practice!

INTRODUCTION TO MLFLOW